Hillsboro Python Machine Learning Meetup

Feb/2017

Ernest Bonat, Ph.D.

Senior Software Engineer

Senior Data Scientist

- 6:00 – 6:40 pm: Pizza, **water only** and networking.
- 6:40 – 6:45 pm: Welcome message by Ernest Bonat, Ph.D.
- 6:45 – 8:00 pm: Presentation and open discussions.
- 8.00 pm – 9.00 pm: Coding and learning session. Bring your Python development laptop!

Why did I create this meetup?

1. Bad traffic to Portland downtown.
2. Hard to find a parking.
3. Bad Python presentation code.
4. No time at all to review the presentation and learn something after the meeting.

We need your support:

1. Need 2 Senior Python Developers for presentation and code review every month (Co-organizers, 4-6 hours a month).

2. Every meeting cost about $200. We need companies to sponsor our meetings.

3. Email Ernest at ebonat@15itresources.com

Our Meetup Mission:

1. *"Come, Listen, Code and Learn".*

2. Finding and presenting best practices of Machine Learning using Python Data Stack.

3. Create great networking place for Hillsboro-Beaverton Data Scientists.

Today Presentation

"Using Python Pandas Library for Data Manipulation and Cleansing"

pandas - an open source library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language. (http://pandas.pydata.org)

Release: 0.19.2 - December 2016

Cheat Sheet

https://github.com/pandas-dev/pandas/blob/master/doc/cheatsheet/Pandas_Cheat_Sheet.pdf

PDF Documentation File

http://pandas.pydata.org/pandas-docs/stable/pandas.pdf

Created by: Wes McKinney, now maintained by Jeff

Reback and many others.

O'Reilly Media Book: “Python for Data Analysis.”

Data Wrangling with Pandas, NumPy, and IPython.

Pandas Cheat Sheet - Python for Data Science

https://www.dataquest.io/blog/pandas-cheat-sheet/

Two main imports:

```
import pandas as pd
import numpy as np
```

Definition

df	Any pandas DataFrame object
s	Any pandas Series object

Importing Data

pd.read_csv(filename)	From a CSV file
pd.read_table(filename)	From a delimited text file (like TSV)
pd.read_excel(filename)	From an Excel file
pd.read_sql(query, connection_object)	Read from a SQL table/database
pd.read_json(json_string)	Read from a JSON formatted string, URL or file.
pd.read_html(url)	Parses an html URL, string or file and extracts tables to a list of dataframes
pd.read_clipboard()	Takes the contents of your clipboard and passes it to read_table()

pd.DataFrame(dict)	From a dict, keys for columns names, values for data as lists

Exporting Data

df.to_csv(filename)	Write to a CSV file
df.to_excel(filename)	Write to an Excel file
df.to_sql(table_name, connection_object)	Write to a SQL table
df.to_json(filename)	Write to a file in JSON format

Create Test Objects

pd.DataFrame(np.random.rand(20,5))	5 columns and 20 rows of random floats
pd.Series(my_list)	Create a series from an iterable my_list
df.index = pd.date_range('1900/1/30', periods=df.shape[0])	Add a date index

Viewing/Inspecting Data

df.head(n)	First n rows of the DataFrame
df.tail(n)	Last n rows of the DataFrame
df.shape()	Number of rows and columns
df.info()	Index, Datatype and Memory information
df.describe()	Summary statistics for numerical columns
s.value_counts(dropna=False)	View unique values and counts

df.apply(pd.Series.value_counts)	Unique values and counts for all columns

Selection

df[col]	Return column with label col as Series
df[[col1, col2]]	Return Columns as a new DataFrame
s.iloc[0]	Selection by position
s.loc['index_one']	Selection by index
df.iloc[0,:]	First row
df.iloc[0,0]	First element of first column

Data Cleaning

df.columns = ['a','b','c']	Rename columns
pd.isnull()	Checks for null Values, Returns Boolean Arrray
pd.notnull()	Opposite of pd.isnull()
df.dropna()	Drop all rows that contain null values
df.dropna(axis=1)	Drop all columns that contain null values
df.dropna(axis=1,thresh=n)	Drop all rows have have less than n non null values
df.fillna(x)	Replace all null values with x
s.fillna(s.mean())	Replace all null values with the mean (mean can be replaced with almost any

	function from the statistics section)
s.astype(float)	Convert the datatype of the series to float
s.replace(1,'one')	Replace all values equal to 1 with 'one'
s.replace([1,3],['one','three'])	Replace all 1 with 'one' and 3 with 'three'
df.rename(columns=lambda x: x + 1)	Mass renaming of columns
df.rename(columns={'old_name': 'new_ name'})	Selective renaming
df.set_index('column_one')	Change the index
df.rename(index=lambda x: x + 1)	Mass renaming of index

Join/Comine

df1.append(df2)	Add the rows in df1 to the end of df2 (columns should be identical)
df.concat([df1, df2],axis=1)	Add the columns in df1 to the end of df2 (rows should be identical)
df1.join(df2,on=col1,how='inner')	SQL-style join the columns in df1 with the columns on df2 where the rows for col have identical values. how can be one of 'left', 'right', 'outer', 'inner'

Descriptive Statistics

(These can all be applied to a series as well)

df.describe()	Summary statistics for numerical columns
df.mean()	Return the mean of all columns
df.corr()	Finds the correlation between columns in a DataFrame
df.count()	Counts the number of non-null values in each DataFrame column
df.max()	Finds the highest value in each column
df.min()	Finds the lowest value in each column
df.median()	Finds the median of each column

df.std()	Finds the standard deviation of each column

Why pandas?

- Heterogeneous data types
- Easy, fast missing data handling
- Easier to write generic code
- Labeled data (numpy mostly assumes index == label)
- Relational data

pandas Data Structures Objects

1. Series
2. DataFrame
3. Panel

Series

A one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). The axis labels are collectively referred to as the **index**.

```
s = pd.Series(data, index=index)
```

Where: **data** can be: Python dictionary, ndarray (n-dimensional array or any scalar value (like 10)

Example:

```
s = pd.Series(np.random.randn(5))
print(s)
```

Result:

```
0    0.3674
1   -0.8230
2   -1.0295
3   -1.0523
4   -0.8502
dtype: float64
```

DataFrame

A 2-dimensional labeled data structure with rows and columns of potentially different types (similar to Microsoft Excel spreadsheet or SQL database table)

```
df = pd.DataFrame(data, …)
```

DataFrame accepts many different kinds of input:

- Dictonary of 1D ndarrays, lists, dicts, or Series
- 2-D numpy.ndarray
- Structured or record ndarray
- A Series
- Another DataFrame

Example:

```
dictionary = {“one” : [1., 2., 3., 4.], “two” : [4., 3., 2., 1.]}
df = pd.DataFrame(dictionary)
```

Result:

```
   one  two
0  1.0  4.0
1  2.0  3.0
2  3.0  2.0
3  4.0  1.0
```

Indexing / Selection

The basics of indexing are as follows:

Operation	Syntax	Result
Select column	df[col]	Series
Select row by label	df.loc[label]	Series
Select row by integer location	df.iloc[loc]	Series
Slice rows	df[5:10]	DataFrame
Select rows by boolean vector	df[bool_vec]	DataFrame

Panel

A 3-dimensional labeled data structure. It's less-used today!

Missing Data

Missing Data is define as Non-available (NA), null or "not present for whatever reason"

pandas uses "NaN" (Non-a-Number) or "nan" internally for simplicity and performance reasons

In CSV file:

one	two	three	four	five	timestamp
	2.1	3.1	bar	1	
	2.3	3.2		0	1/1/2017
1.3		3.3	bar	1	2/1/2017
1.4	2.4		bar		3/1/2017
1.5	2.5	3.5	bar	0	

In pandas DataFrame:

..	one	two	three	four	five	timestamp
0	**nan**	2.1	3.1	bar	1	nan
1	**nan**	2.3	3.2	**nan**	0	1/1/2017
2	1.3	**nan**	3.3	bar	1	2/1/2017
3	1.4	2.4	nan	bar	**nan**	3/1/2017
4	1.5	2.5	3.5	bar	0	**nan**

Data Science Two Main Tasks:

1	**Data Cleansing**	**60% - 70% work**
2	Data Analytics	40% - 30% work

Data Cleansing very important task. Be careful with **“Garbage IN – Garbage OUT”**

Beginning Steps:

1. Organize Input and Output Data Files Path Name
2. Import Data File to pandas DataFrame
3. Get Number of Rows and Columns
4. Get Index, Datatype and Memory Information
5. Remove Duplicates Rows
6. Fill Nan Values (Mean, Median, Defaults, etc.)

7. Remove Rows by Row/Column Conditions
8. Replace Values by Row/Column Conditions

Coding session:

Provide the data cleansing for “bad_titanic.csv” file.
Create a good_titanic.csv file following the requirements:

- Remove duplicates records and keep first found
- Survived should be 0 or 1 only
- Class should be 1, 2 and 3 only
- Sex should be “male” or “female” only
- Empty Age number replace with “NA”
- SibSp should be <= 8
- Parch should be <= 6
- Empty Cabin replace with “NA”
- Embarked should be S, C and Q only

Presentation source code:

(https://github.com/ebonat/hillsboro_machine_learning_02_2017)